

# 强化学习及其应用

Reinforcement Learning and Its Applications

## 第八章 蒙特卡罗树搜索

Monte Carlo Tree Search

授课人：周晓飞

zhouxiaofei@iie.ac.cn

2023-6-15

# 第八章 蒙特卡罗树搜索

## 8.1 Decision Time Planning

## 8.2 蒙特卡罗树（MCTS）搜索

## 8.3 AlphaGo

### 8.3.1 Policy Gradient

### 8.3.2 MCTS 方法

# 第八章 蒙特卡罗树搜索

## 8.1 Decision Time Planning

## 8.2 蒙特卡罗树（MCTS）搜索

## 8.3 AlphaGo

### 8.3.1 Policy Gradient

### 8.3.2 MCTS 方法

# Decision Time Planning

■ **两种 Planning** (都是基于已有 episodes 不充分的假设, 需要生成 episodes)

■ **Background Planning**

目标: 学习背景和环境的 Model, 进而采集 Model 的 episodes 来优化策略;

典型的方法: Dynamic Programming 和 Dyna。

■ **Decision-Time Planning**

目标: 不以学习背景环境为目的, 重点要学习状态  $S_t$  下的最优策略  $a_t$ ; 结合了策略改善和环境模拟。


典型的方法: MCTS

# Decision Time Planning

## ■ 例如 MTCS: MC 估计+启发式搜索

Episodes 本质源于一颗以  $S_t$  为 root 的树，MC 估计  $S_t$  的值函数；  
通过启发式搜索来引导对树上 episodes 的采集。

值估计的随机逼近: 优先采集  $Q(s,a)$  大的行动，使得  $V(s)$  不断改善



No need to solve whole MDP, just sub-MDP starting from **now**

启发式搜索原则: 优先采集  $Q(s,a)$  大的行动分支

# Decision Time Planning

## 问题描述

- Simulate episodes of experience from now with the model

$$\{s_t^k, a_t^k, r_{t+1}^k, \dots, s_T^k\}_{k=1}^K \sim \mathcal{M}_\nu$$

- Apply model-free RL to simulated episodes
  - Monte-Carlo control  $\rightarrow$  Monte-Carlo search
  - Sarsa  $\rightarrow$  TD search

# 第八章 蒙特卡罗树搜索

8.1 Decision Time Planning

**8.2 蒙特卡罗树（MCTS）搜索**

8.3 AlphaGo

8.3.1 Policy Gradient

8.3.2 MCTS 方法

## Simple Monte-Carlo Search

- Given a model  $\mathcal{M}_\nu$  and a policy  $\pi$
- For each action  $a \in \mathcal{A}$ 
  - Simulate  $K$  episodes from current (real) state  $s_t$

$$\{s_t, a, r_{t+1}^k, s_{t+1}^k, a_{t+1}^k, \dots, s_T^k\}_{k=1}^K \sim \mathcal{M}_\nu, \pi$$

- Evaluate actions by mean return (Monte-Carlo evaluation)

$$Q(s_t, a) = \frac{1}{K} \sum_{k=1}^K v_t \xrightarrow{P} Q^\pi(s_t, a)$$

- Select current (real) action with maximum value

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q(s_t, a)$$



# 蒙特卡罗树搜索

## Monte-Carlo Tree Search

- Given a model  $\mathcal{M}_\nu$

- Simulate  $K$  episodes from current state  $s_t$  using current simulation policy  $\pi$

$$\{s_t, a_t^k, r_{t+1}^k, s_{t+1}^k, \dots, s_T^k\}_{k=1}^K \sim \mathcal{M}_\nu, \pi$$

- Build a search tree containing visited states and actions

- Evaluate states  $Q(s, a)$  by mean return of episodes from  $s, a$

$$Q(s, a) = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{u=t}^T \mathbf{1}(s_u, a_u = s, a) v_u \xrightarrow{P} Q^\pi(s, a)$$

- After search is finished, select current (real) action with maximum value in search tree

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q(s_t, a)$$

- In MCTS, the simulation policy  $\pi$  improves
- Each simulation consists of two phases (in-tree, out-of-tree)
  - Tree policy (improves): pick actions to maximise  $Q(s, a)$
  - Default policy (fixed): pick actions randomly
- Repeat (each simulation)
  - Evaluate states  $Q(s, a)$  by Monte-Carlo evaluation
  - Improve tree policy, e.g. by  $\epsilon$ -greedy( $Q$ )

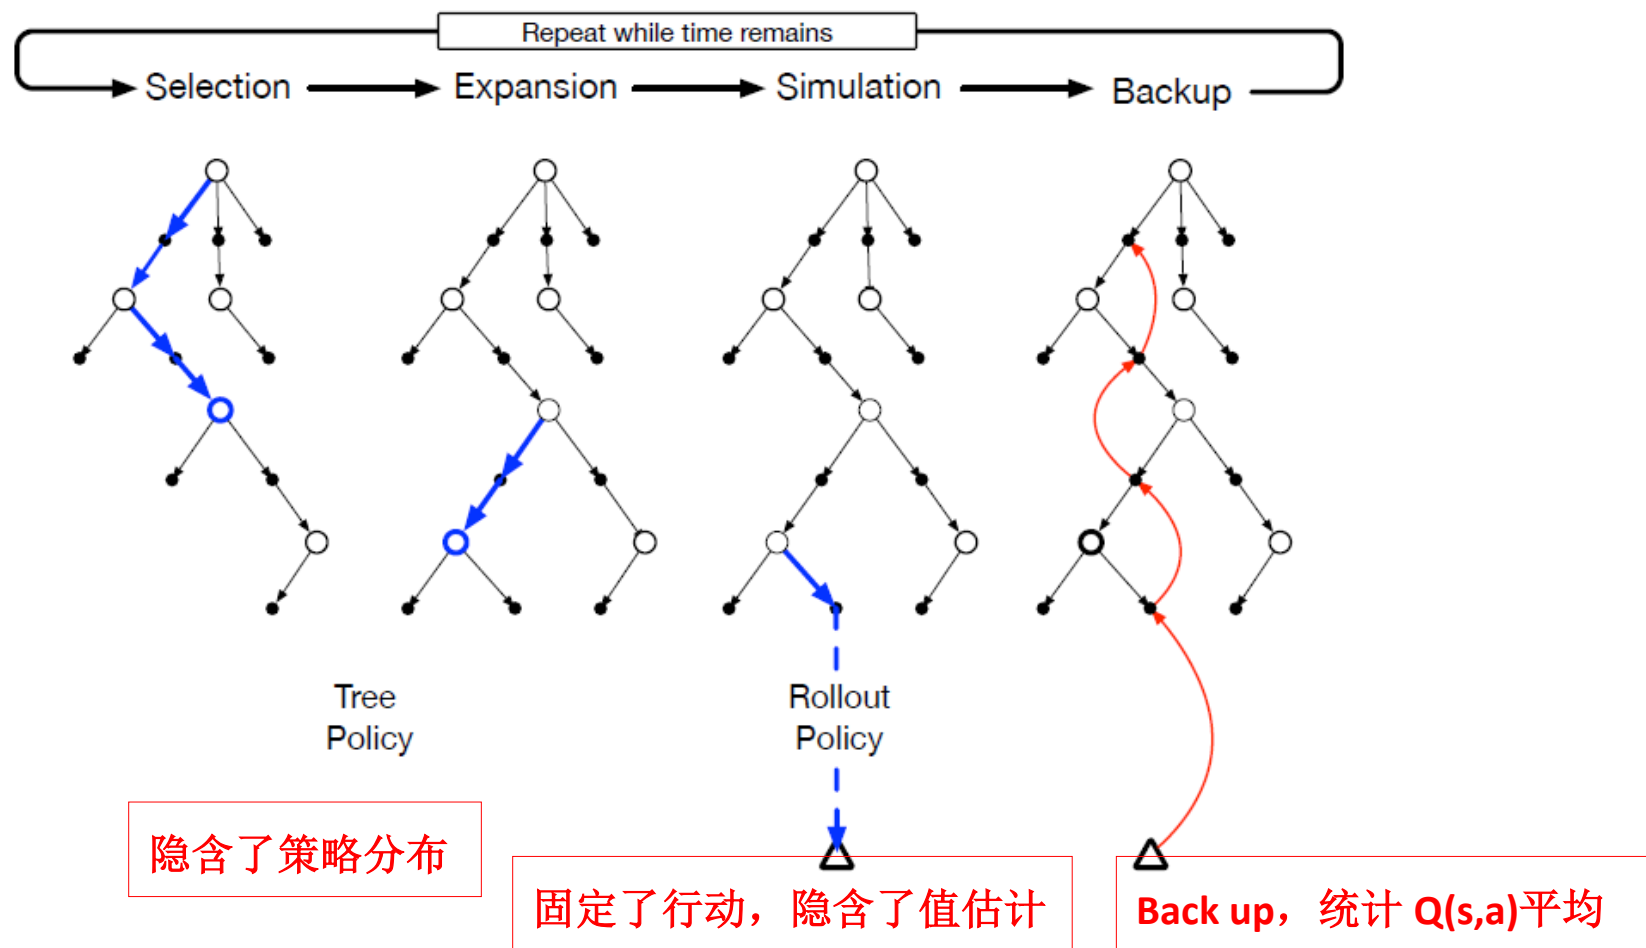
# 蒙特卡罗树搜索

## Monte-Carlo Tree Search

- In MCTS, the simulation policy  $\pi$  **improves**
- Each simulation consists of two phases (in-tree, out-of-tree)
  - **Tree policy** (improves): pick actions to maximise  $Q(s, a)$
  - **Default policy** (fixed): pick actions randomly
- Repeat (each simulation)
  - **Evaluate** states  $Q(s, a)$  by Monte-Carlo evaluation
  - **Improve** tree policy, e.g. by  $\epsilon$  – greedy( $Q$ )
- **Monte-Carlo control** applied to **simulated experience**
- Converges on the optimal search tree,  $Q(s, a) \rightarrow Q^*(s, a)$

# 蒙特卡罗树搜索

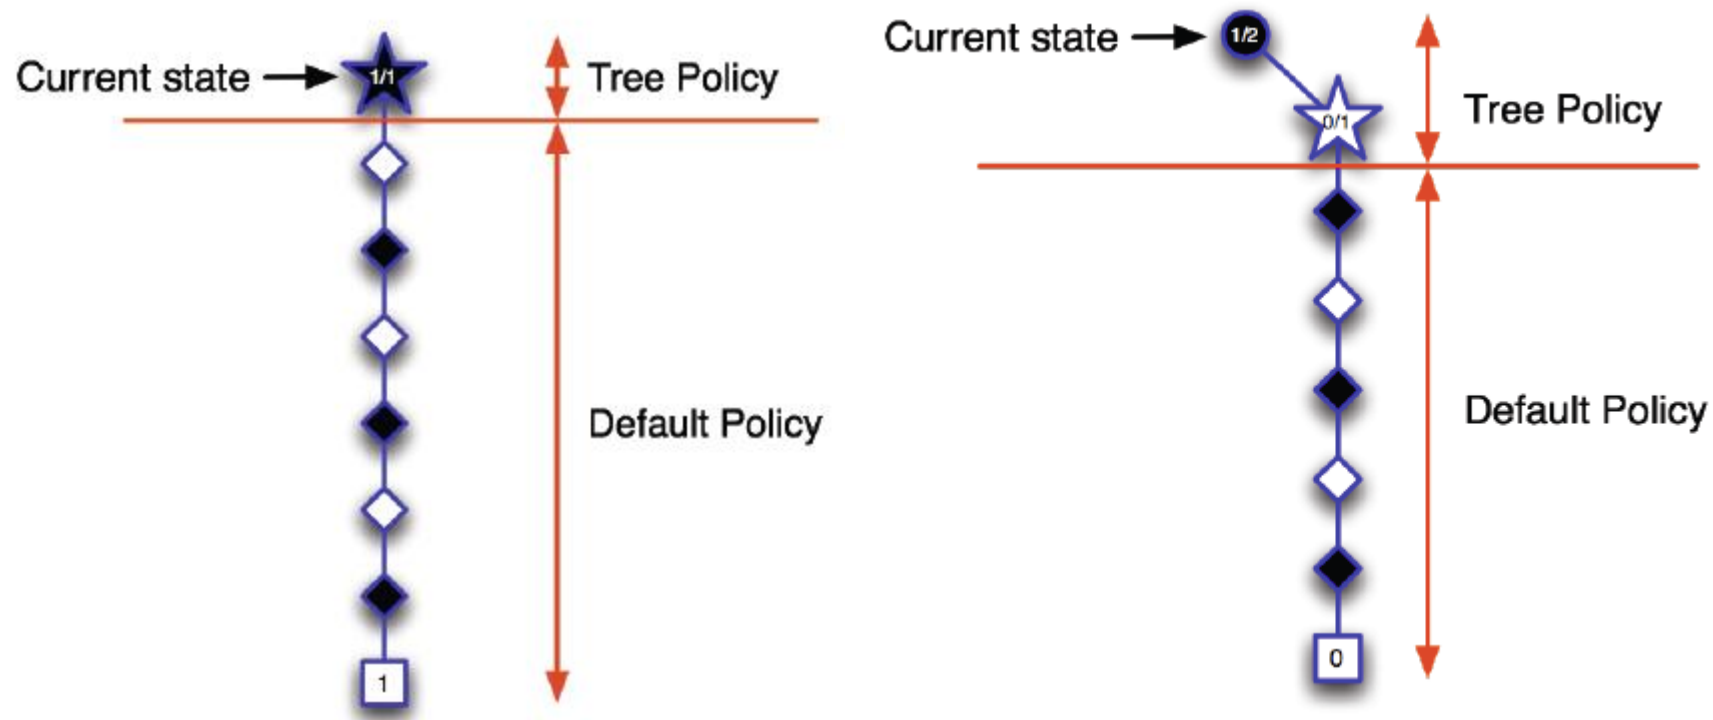
## Basic Version of MCTS





## Basic Version of MCTS

1. **Selection.** Starting at the root node, a *tree policy* based on the action values attached to the edges of the tree traverses the tree to select a leaf node.
2. **Expansion.** On some iterations (depending on details of the application), the tree is expanded from the selected leaf node by adding one or more child nodes reached from the selected node via unexplored actions.
3. **Simulation.** From the selected node, or from one of its newly-added child nodes (if any), simulation of a complete episode is run with actions selected by the rollout policy. The result is a Monte Carlo trial with actions selected first by the tree policy and beyond the tree by the rollout policy.
4. **Backup.** The return generated by the simulated episode is backed up to update, or to initialize, the action values attached to the edges of the tree traversed by the tree policy in this iteration of MCTS. No values are saved for the states and actions visited by the rollout policy beyond the tree. Figure 8.13 illustrates this by showing a backup from the terminal state of the simulated trajectory directly to the state–action node in the tree where the rollout policy began (though in general, the entire return over the simulated trajectory is backed up to this state–action node).


# 蒙特卡罗树搜索




# 蒙特卡罗树搜索



# 蒙特卡罗树搜索



# 蒙特卡罗树搜索





# 蒙特卡罗树搜索

## MCTS Algorithm

### Algorithm 1: General MCTS approach

```
function MCTSSEARCH( $s_0$ )  
  create root node  $v_0$  with state  $s_0$   
  while within computational budget do  
     $v_l \leftarrow \text{TREEPOLICY}(v_0)$   
     $\Delta \leftarrow \text{DEFAULTPOLICY}(s(v_l))$   
    BACKUP( $v_l, \Delta$ )  
  return  $a(\text{BESTCHILD}(v_0))$ 
```

# 蒙特卡罗树搜索

## MCTS Algorithm

### Algorithm 2: The UCT algorithm

```
function UCTSEARCH( $s_0$ )
  create root node  $v_0$  with state  $s_0$ 
  while within computational budget do
     $v_l \leftarrow \text{TREEPOLICY}(v_0)$ 
     $\Delta \leftarrow \text{DEFAULTPOLICY}(s(v_l))$ 
    BACKUP( $v_l, \Delta$ )
  return  $a(\text{BESTCHILD}(v_0, 0))$ 

function TREEPOLICY( $v$ )
  while  $v$  is nonterminal do
    if  $v$  not fully expanded then
      return EXPAND( $v$ )
    else
       $v \leftarrow \text{BESTCHILD}(v, Cp)$ 
  return  $v$ 
```

```
function EXPAND( $v$ )
```

```
  choose  $a \in$  untried actions from  $A(s(v))$ 
  add a new child  $v'$  to  $v$ 
    with  $s(v') = f(s(v), a)$ 
    and  $a(v') = a$ 
  return  $v'$ 
```

```
function BESTCHILD( $v, c$ )
```

```
  return  $\arg \max_{v' \in \text{children of } v} \frac{Q(v')}{N(v')} + c \sqrt{\frac{2 \ln N(v)}{N(v)}}$ 
```

```
function DEFAULTPOLICY( $s$ )
```

```
  while  $s$  is non-terminal do
    choose  $a \in A(s)$  uniformly at random
     $s \leftarrow f(s, a)$ 
  return reward for state  $s$ 
```

```
function BACKUP( $v, \Delta$ )
```

```
  while  $v$  is not null do
     $N(v) \leftarrow N(v) + 1$ 
     $Q(v) \leftarrow Q(v) + \Delta(v, p)$ 
     $v \leftarrow \text{parent of } v$ 
```

# 蒙特卡罗树搜索

## MCTS 优点

- Highly selective best-first search
- Evaluates states *dynamically* (unlike e.g. DP)
- Uses sampling to break curse of dimensionality
- Works for “black-box” models (only requires samples)
- Computationally efficient, anytime, parallelisable

# 蒙特卡罗树搜索

## 分析:

```
function BESTCHILD( $v, c$ )  
return  $\arg \max_{v' \in \text{children of } v} \frac{Q(v')}{N(v')} + c \sqrt{\frac{2 \ln N(v)}{N(v)}}$ 
```

较大 $Q(s_t, a)$ 值的 $a$ 分支被搜索次数多, 对应的 $N(s_t, a)$ 较大, 搜索频率逼近树策略分布.

$$\begin{aligned} V(s) &= \sum_a \pi(a|s) Q(s, a) \\ &= \frac{1}{N(s)} \sum_a N(s, a) Q(s, a) \end{aligned}$$

# 第八章 蒙特卡罗树搜索

8.1 Decision Time Planning

8.2 蒙特卡罗树搜索

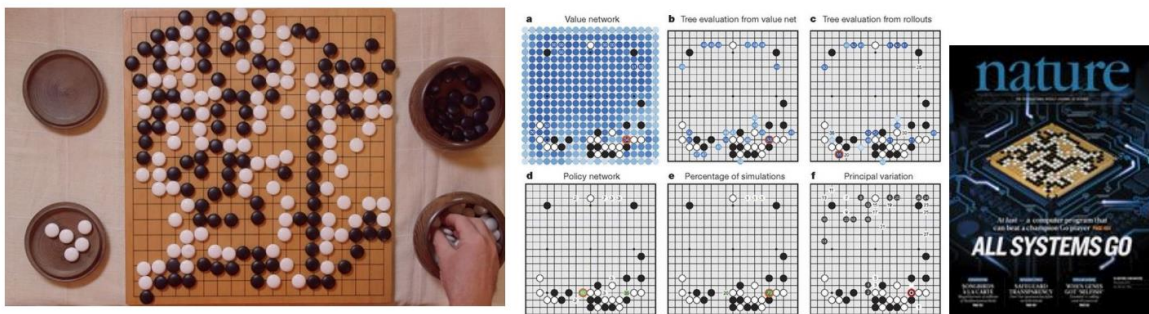
**8.3 AlphaGo**

8.3.1 Policy Gradient

8.3.2 MCTS 方法

# AlphaGo

## 围棋问题



- S: 棋局;
- A: 落子动作;
- Model:

状态转移: 互搏弈导致的状态转移,  $S \rightarrow S'$ , 没有直观的  $P_{ss'}$  ;

$(S, a_+) \rightarrow S_+$ , (正方落子后进入状态  $S_+$ );  $(S_+, a_-) \rightarrow S'$  (反方落子后进入状态  $S'$ )

回报值: 最终是否获胜, 获胜为 1, 否则为-1; 需要完整的 episode.

问题: 在已有棋谱或专家经验 (training episodes) 作为先验, 当前棋局  $S_t$  的最优行动  $a_t$  ?

## Abstract

### ■ 提出两个方法:

#### 1. 不采用前向 lookahead Search, 学习策略网络 (Policy Network)

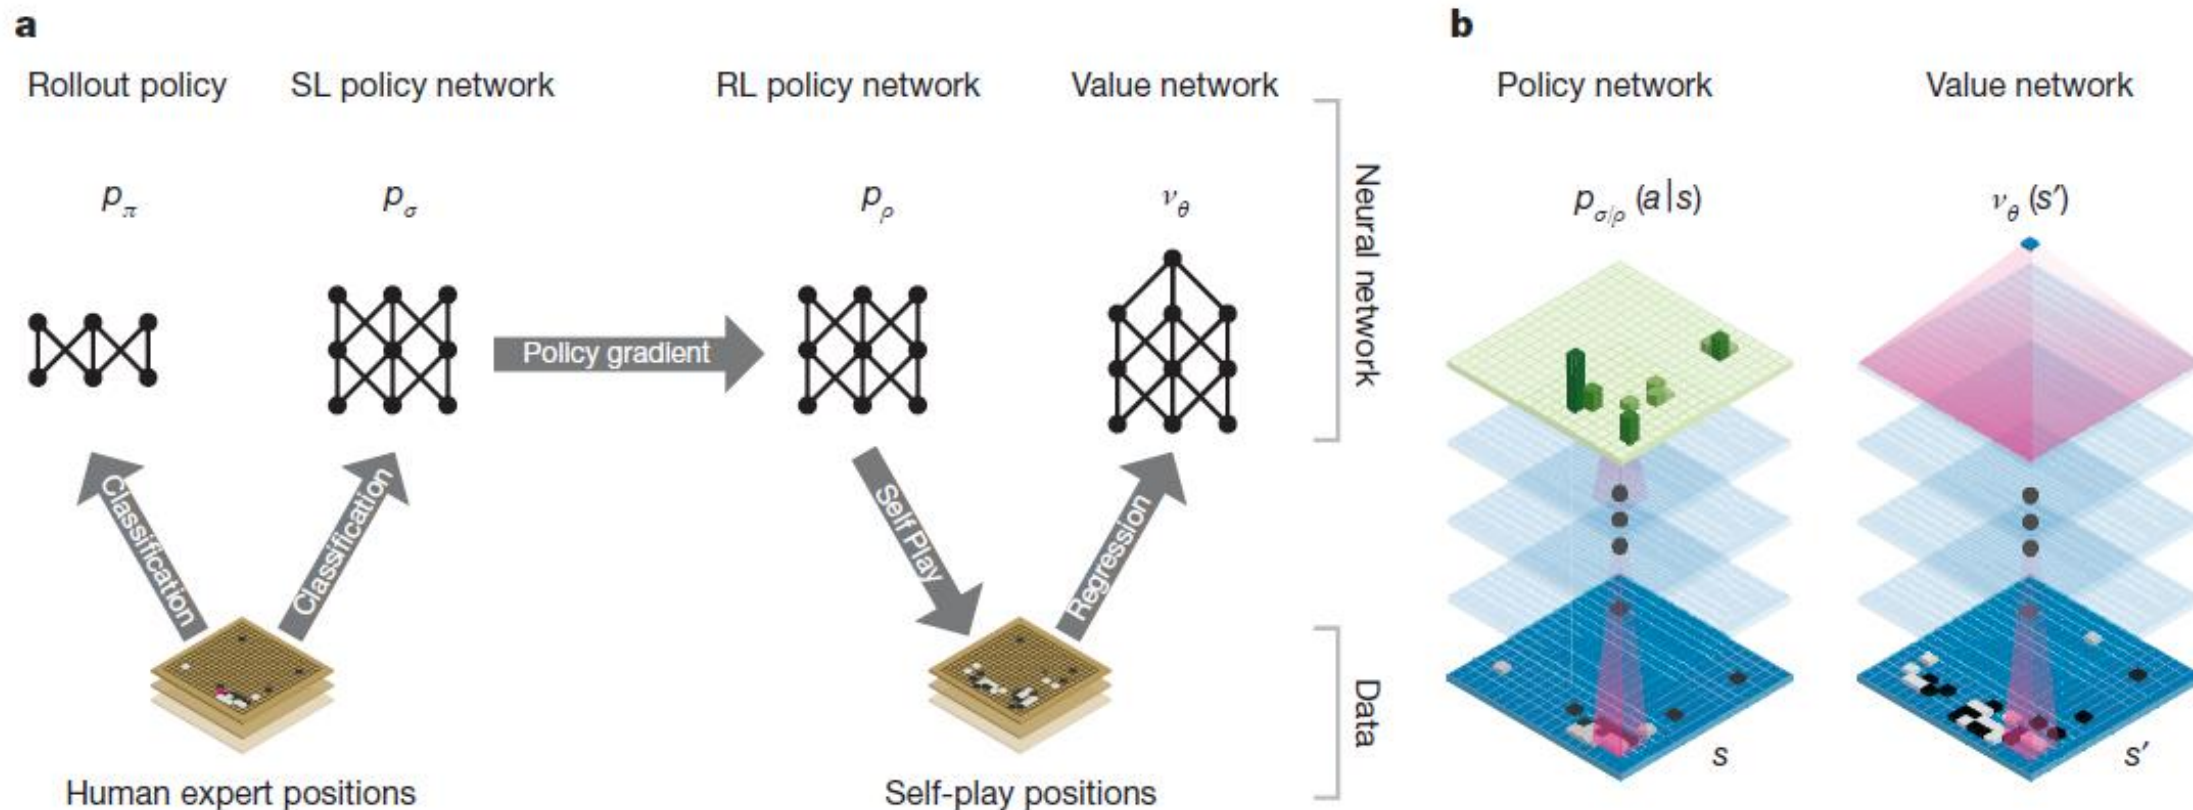
结合 value networks 和 policy network: Value networks 估计棋局的值函数,  
Policy networks 选择行动  $a_t$

#### 2. 采用前向 lookahead Search 的 MTCS 方法

结合 value networks, policy network 和 MTCS

两种方法不同: 前者学习 Policy probability 决定  $a_t$ , 后者是搜最优的  $a_t$

# AlphaGo





## 方法一: Policy Gradient

- 目标: 学习一个策略网络  $p_{\rho}(a_t | s_t)$
- 参数更新采用策略梯度方法:

$$\Delta \rho \propto \frac{\partial \log p_{\rho}(a_t | s_t)}{\partial \rho} z_t$$

初始化网络:

专家经验训练了一个先验策略网络(supervised learning (SL) policy network)作为初始化

$$\Delta \sigma \propto \frac{\partial \log p_{\sigma}(a | s)}{\partial \sigma}$$

## 方法一: Policy Gradient

### ■ 方法流程

#### 1. 学习 SL Policy Network

sampled state-action pairs  $(s, a)$ , using stochastic gradient ascent to maximize the likelihood of the human move  $a$  selected in state  $s$

$$\Delta\sigma \propto \frac{\partial \log p_{\sigma}(a|s)}{\partial \sigma}$$

- 13-layer policy network, which we call the SL policy network, from 30 million positions from the KGS Go Server.
- Data: randomly sampled state-action pairs  $(s,a)$  on human moves.

## 方法一: Policy Gradient

### 2. 学习 RL Policy Network $p_\rho(a_t|s_t)$

$$\Delta\rho \propto \frac{\partial \log p_\rho(a_t|s_t)}{\partial \rho} z_t$$

- Its weights  $\rho$  are initialized to the same values,  $\rho = \sigma$ .
- Episodes are from playing games between the **current policy network  $p_\rho$**  and a **randomly selected previous iteration of the policy network**.
- Outcomes  $z_t$ : +1 for winning and -1 for losing.

## 方法一: Policy Gradient


### 3. 学习 Value Network $V_{\theta}(S)$

目的: 值函数  $V_{\theta}(S)$  替代  $z_t$

$$\Delta\theta \propto \frac{\partial v_{\theta}(s)}{\partial \theta} (z - v_{\theta}(s))$$

- Train data: we generated a new self-play data set consisting of 30 million distinct positions, each sampled from a separate game. Each game was played between the RL policy network and itself until the game terminated

## 方法二：MCTS



以  $S_t$  为 root 的子树不断被扩展，  
优先搜索和扩展回报值大的节点，  
即  $Q(s,a)$  大的节点。

## 方法二: MCTS

### AlphaGo with MCTS

#### 1. Selection

- At each of these time steps,  $t < L$ , **an action** is selected according to the statistics in the search tree:

$$a_t = \operatorname{argmax}_a (Q(s_t, a) + u(s_t, a))$$

$$u(s, a) = c_{\text{puct}} P(s, a) \frac{\sqrt{\sum_b N_r(s, b)}}{1 + N_r(s, a)}$$

where  $c_{\text{puct}}$  is a constant determining the level of exploration; this search control strategy initially prefers actions with high prior probability and low visit count, but asymptotically prefers actions with high action value.

选择的 $a$ 可能不存在后继结点，是否扩展后继结点依据Expansion。

## 方法二: MCTS

### AlphaGo with MCTS

#### 树结构

Each node  $s$  in the search tree contains edges  $(s, a)$  for all legal actions  $a \in \mathcal{A}(s)$ .  
Each edge stores a set of statistics,

$$\{P(s, a), N_v(s, a), N_r(s, a), W_v(s, a), W_r(s, a), Q(s, a)\}$$

where  $P(s, a)$  is the prior probability,  $W_v(s, a)$  and  $W_r(s, a)$  are Monte Carlo estimates of total action value, accumulated over  $N_v(s, a)$  and  $N_r(s, a)$  leaf evaluations and rollout rewards, respectively, and  $Q(s, a)$  is the combined mean action value for that edge. Multiple simulations are executed in parallel on separate search threads.

## 方法二: MCTS

### AlphaGo with MCTS

#### 2. Expansion

- When the **visit (selected) count exceeds a threshold**,  $N_r(s, a) > n_{thr}$ , the successor state  $s' = f(s, a)$  is added to the search tree. The new node is initialized to:

$$\{N(s', a) = N_r(s', a) = 0, W(s', a) = W_r(s', a) = 0, P(s', a) = p_\sigma(a|s')\}.$$



## 方法二: MCTS

### ■ AlphaGo with MCTS

#### 3. Evaluation

- These evaluations are combined, using a mixing parameter  $\lambda$ , into a leaf evaluation  $V(S_L)$ ;

$$V(s_L) = (1 - \lambda)v_{\theta}(s_L) + \lambda z_L$$

- The leaf node is evaluated in two very different ways:
  - (1) First, by the value network  $V_{\theta}(S_L)$ ;
  - (2) Second, by the outcome  $Z_L$  of a random rollout played out until terminal step  $T$  using the fast rollout policy  $p_{\pi_i}$ ;

## 方法二: MCTS


### AlphaGo with MCTS

#### 4. Backup

$$N(s, a) = \sum_{i=1}^n 1(s, a, i)$$
$$Q(s, a) = \frac{1}{N(s, a)} \sum_{i=1}^n 1(s, a, i) V(s_L^i)$$

where  $s_L^i$  is the leaf node from the  $i$ th simulation, and  $1(s, a, i)$  indicates whether an edge  $(s, a)$  was traversed during the  $i$ th simulation. Once the search is complete, the algorithm chooses the most visited move from the root position.

# AlphaGo



# 本讲参考文献

1. Richard S. Sutton and Andrew G. Barto. Reinforcement Learning: An Introduction. (Second edition, in progress, draft).
2. David Silver, Slides@ 《Reinforcement Learning: An Introduction》, 2016.
3. Browne C.B., Daniel W. Cowling P.I, and Samothrakis S. A Survey of Monte Carlo Tree Search Methods, IEEE Transactions on Computational Intelligence and AI In Games, vol. 4, no. 1.
4. David Silver, and Aja Huang et.al. Mastering the game of Go with deep neural networks and tree search. Nature 2016.